

KNIGHTEC

ARTIFICIAL INTELLIGENCE AND MACHINE LEARNING

Welcome

18.00

Grab some food and drink

18.20

Region Manager Carl Falk – Introduction

18.30

AI/ML seminar part 1 - Stefan Danielsson

19.15

Pause and “fika”

19.30

AI/ML seminar part 2 – Jonas Dürango

20.15

Finish

“Game AI has always been used as a measure of progress in AI, throughout it’s history”.

Approaching Machine Learning
with Artificial Intelligence, Gadgets & Games

Stefan Danielsson



Experience

Software Developer, Knightec



MDH - Computer Science

Artificial Intelligence

- Studies
- Thesis EEG signal classification and feature selection

Hobbies

Programming

Robots

2D games

3D printing

CNC milling and turning

How familiar are you with AI?

- a) I have never heard of it. What is this dark thing you speak of?
- b) I am not sure what it is, but I have my guesses
- c) I have read a lot about AI in magazines, articles and on the web
- d) I have studied AI to some extent, and I have some practical experience
- e) I incorporate AI in my work, every day
- f) I probably know more about AI, then anyone else in this room

Artificial Intelligence history in 60 seconds

Thinking machines

In 1950, Alan Turing published a paper, in which he speculated about the possibility of creating machines that can think. This paper lay the foundation for how intelligent behavior could be distinguishable from that of a human.

Birth of Game AI

In 1951, the first renowned checkers program and chess program, was written..

Mimicking problem solving

In 1955, a program called “Logic Theorist” was developed. It was the first program deliberately engineered to mimic the problem solving skills of a human mathematician. Among other things, it proved mathematical theorems, and even developed more elegant proofs for some theorems.

Industrial robots

In 1961, the company Unimation deployed their industrial robot “Unimate”, on an automobile assembly line, for GM.

Expert systems

In 1965, the first “expert systems” started to take root. The systems could answer questions, or solve problems about a specific domain of knowledge, using logical rules that were derived from the knowledge of experts.

Bayesian methods

In the 1960s, major breakthroughs were made for probabilistic inference, in the field of Statistics and Machine learning.

First autonomous vehicle

In 1979, the “Stanford Cart” becomes the first computer-controlled, autonomous vehicle, after it successfully traverses a chair-filled room.

Backgammon AI

In 1979, a program called “BKG” defeats the reigning backgammon world champion.

Common sense knowledge

In middle of the 1980s, the “Cyc” project was introduced. It was the first attempt to solve the common sense knowledge problem, by creating a massive database, that would contain all the ordinary facts, that an average person knows.

Neural networks

In 1986, the concept of backpropagation for neural networks is introduced, and the concepts of neural networks becomes widely accepted and used.

Intelligent agents

In the 1990s, IA became popular and widely used. IA is an autonomous entity, which observes through sensors, and acts upon an environment using actuators, and directs its activity towards achieving goals.

Checkers AI

In 1994 - Checkers world champion lost to a checkers program called “Chinook”.

Vector machines

In 1995, Vector machines started to become widely used.

“No Hands Across America”

In 1995, a semi-autonomous car drove coast-to-coast across the United States with computer-controlled steering.

Othello AI

In 1997, a program called “Logistello” defeated the world champion.

Deep Blue versus Kasparov

In 1997, the computer “Deep Blue” wins over the reigning chess champion Garry Kasparov.

Google Translate

Google launched their human language translator in 2006.

Atari 2600

In 2013, DeepMind successfully could learn to master just about any Atari 2600 games, and beat the world champions in some of them.

AlphaGo

In 2016, Google's AlphaGo program becomes the first program to beat an professional human Go player.

Star Craft

In 2016, DeepMind and Blizzard initiated an AI collaboration towards StarCraft, as a possible future challenge, since it requires a high level of strategic thinking and handling imperfect information. Blizzard develops a Machine Learning API that gives researchers and developers hooks into the game.

OpenAI & Dota II

During “The International” in 2017, OpenAI let a machine-learned bot play 1v1 demonstration games, against professional Dota 2 players. They all got wrecked!

Chess master in 4 hs

This Christmas, after being programmed with only the rules of chess (no strategies), in just four hours, Google’s “AlphaZero” had mastered the game to the extent it was able to best the highest-rated chess-playing program “Stockfish”.

2010... boom!

IoT

More people and things than ever, are connected.

Availability

Everyone can afford it. Anyone can use it.

Big data & VLDB

The ability to store, capture, analyze, search, secure, share, visualize, query, update, and process large amounts of data.

Cloud computing

*Combined computational power, in a dynamic fashion.
Accessible, from anywhere!*

End-node computing

Decrease in space complexity. Increase in data relevance.

Hardware performance

Hardware acceleration, transistor count, GPU computing, and more...

Deep Learning is kicking into gear!

KNIGHTEC



What is Artificial Intelligence?

KNIGHTEC

Weak-, or applied AI

An intellect that follows rules, that humans gave it, in order to achieve some goal. This kind of intellect can be perceived, as “smart”, but in fact, this intellect is not smart in any way, since it does not in any way resembles the intellect that of a human brain. It lacks the fundamental features that is associated with human intelligence. For example, the ability to reason, associate, relate, generalize, or to be creative. The intellect merely autonomously follows the rules given to it.

Strong-, or generalized AI

An intellect that focuses on mimicking the behavior of a human brain, rather than achieving optimal solutions to a problem. Focus lies on the ability to reason, associate, relate, generalize, or to be creative, in the same fashion that of a human brain. In other words it could be argued that in a “Strong AI”, the reasoning itself is more important than the actual solution.

Superintelligence

A sentient and self-aware intellect, that goes beyond and overcomes the limitations of our known physical existence, and that by some definitions also becomes independent of it. A system that is more intelligent than every person on earth, that has ever lived, combined!

Learning through examples

From basic weak AI, to ML and learning.

Illustrated and visualized through robots, games and applications. Focus on the behavior, and the logic behind the algorithms.

- Autonomous robots
- Basic swarm AI
- Optimization and search trees
- Cumulative rewarding
- Complexity and problem construction
- Evolutionary computation and Genetic Algorithms
- Artificial Neural Networks, and supervised learning

Examples will be available in a GIT-repository for you to try out after the event!







KNIGHTEC

Autonomous robots


The 1985s robotic paradigm

- **Sense** – Gather sensor data from the environment
- **Plan** – Interpret the data, and make an *independent* decision based on it
- **Act** – Perform a suitable action, based on the decision



Sonar sensor

Sensor data	Left	Front	Right
Raw time	-	-	8 ms
Servo Angle	-75°	0°	+75°
PWM	255	255	35
Distance	> 100 mm	> 100 mm	17.5 mm
Hit	False	False	True



Infrared reflectance sensor

Sensor data	Left	Right
PWM	255	35
Hit	False	True



KNIGHTEC

We can approximate parallelism with interrupts and timers.

// Main loop

```
Interrupt sensors();  
Interrupt servo();  
  
while(true)  
{  
    switch(decision)  
    {  
        case decision1:  
        case decision2 :  
        case decision3 :  
            ...  
        default:  
    }  
}
```

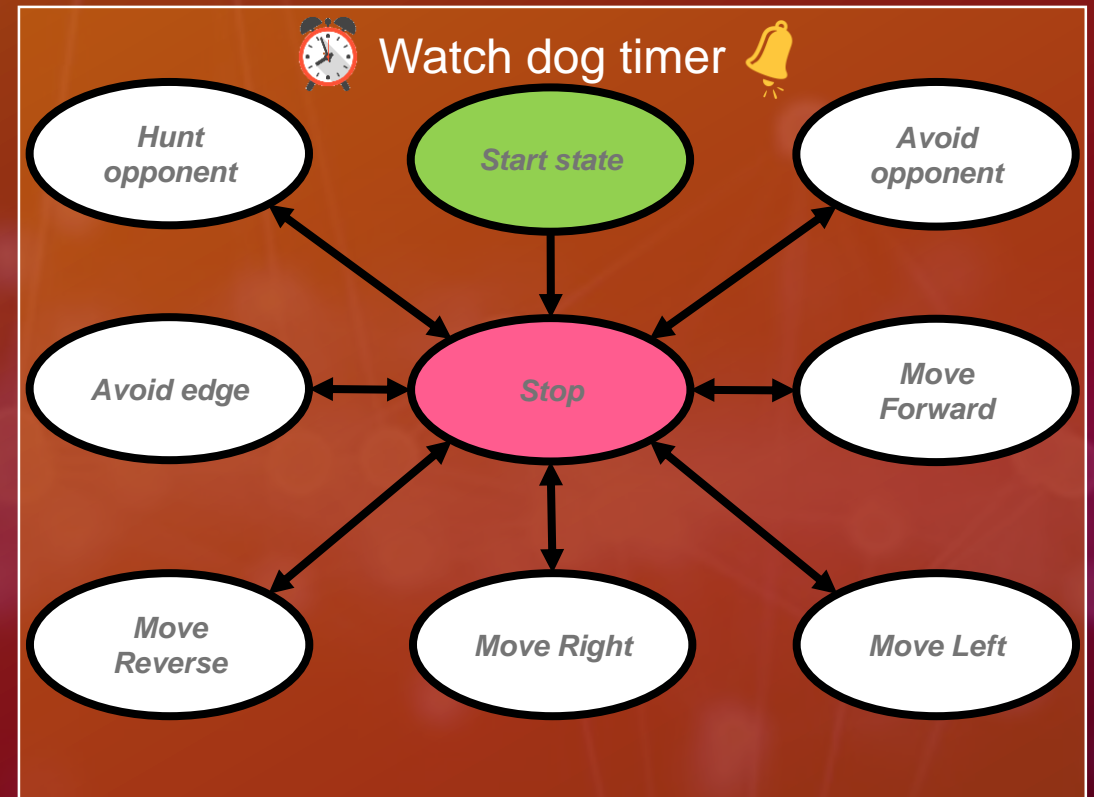
// Sensor(s) interrupt

```
distance = readPing();  
lineHit = readReflectance();
```

// Servo interrupt

```
while (1)  
{  
    sweepServo(direction, step);  
}
```

A simple robot can be thought of as a state machine





Demo

KNIGHTEC

Basic swarm AI

Google AI challenge – AI vs AI

- Ant wars
- Planet wars
- Tron



The Ant wars Engine

Imperfect information – You only know what you can see. You have to explore.

Each turn your bot will be given information for all squares that are visible to your ants.

A bot can issue up to one order for each ant (move one square), during one turn (UP, DOWN, LEFT, RIGHT).



Food = MORE ants!

Ants

Hill (or spawning hole)


Food

KNIGHTEC

The Ant wars Engine

The game then goes through 5 phases, in the following order:


- (1) move all ants (ants that collide in the same square are killed)
- (2) attack enemy ants if within range
- (3) raze ant hills with enemy ants positioned on them
- (4) spawn more ants at hills, and food on map
- (5) gather food next to ants



Demo

KNIGHTEC

Construction and problem representation



"The knapsack problem"

We have one knapsack, and some items.

Items have:

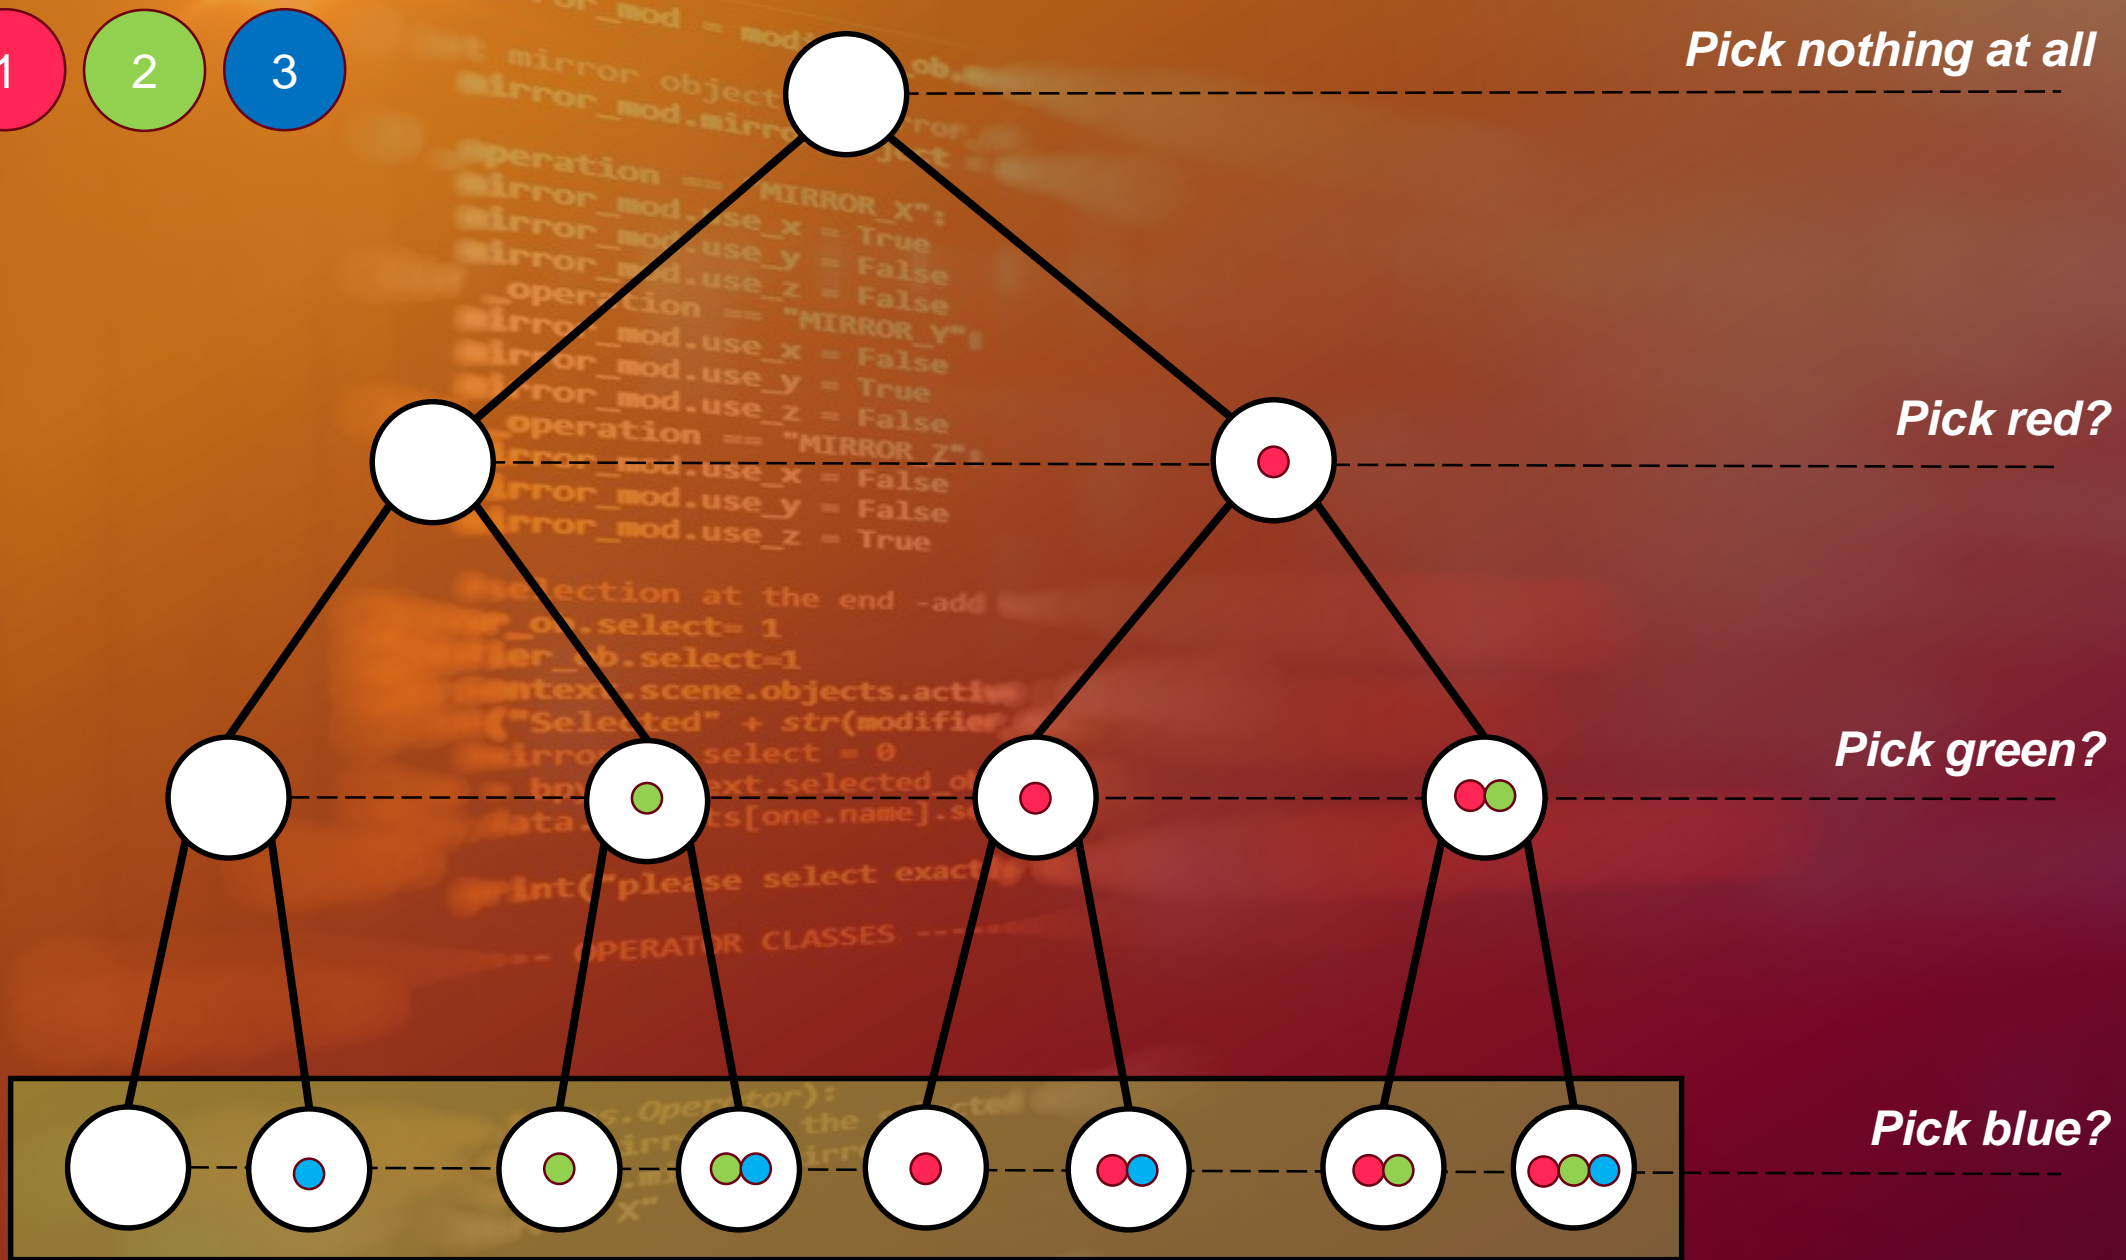
- Weight
- Value

Knapsack has:

- Weight limit

The problem:

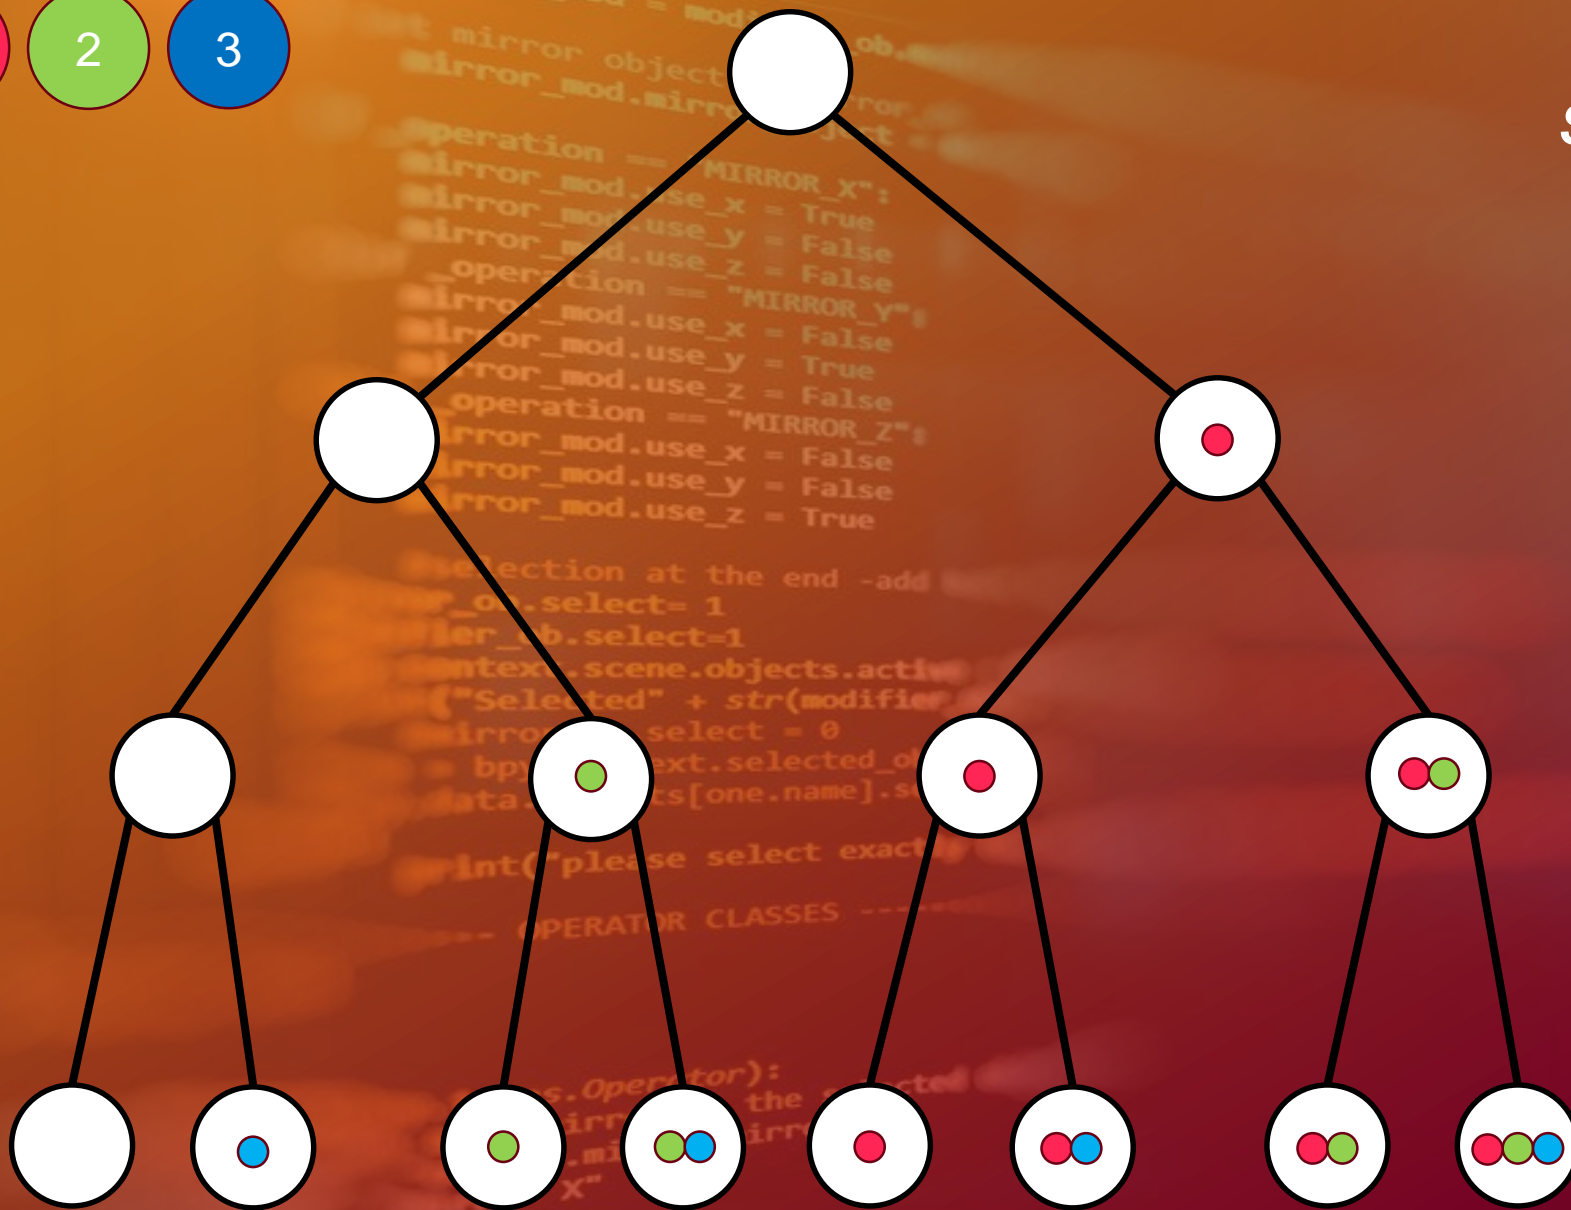
What items do we put in the knapsack, for maximum total value?

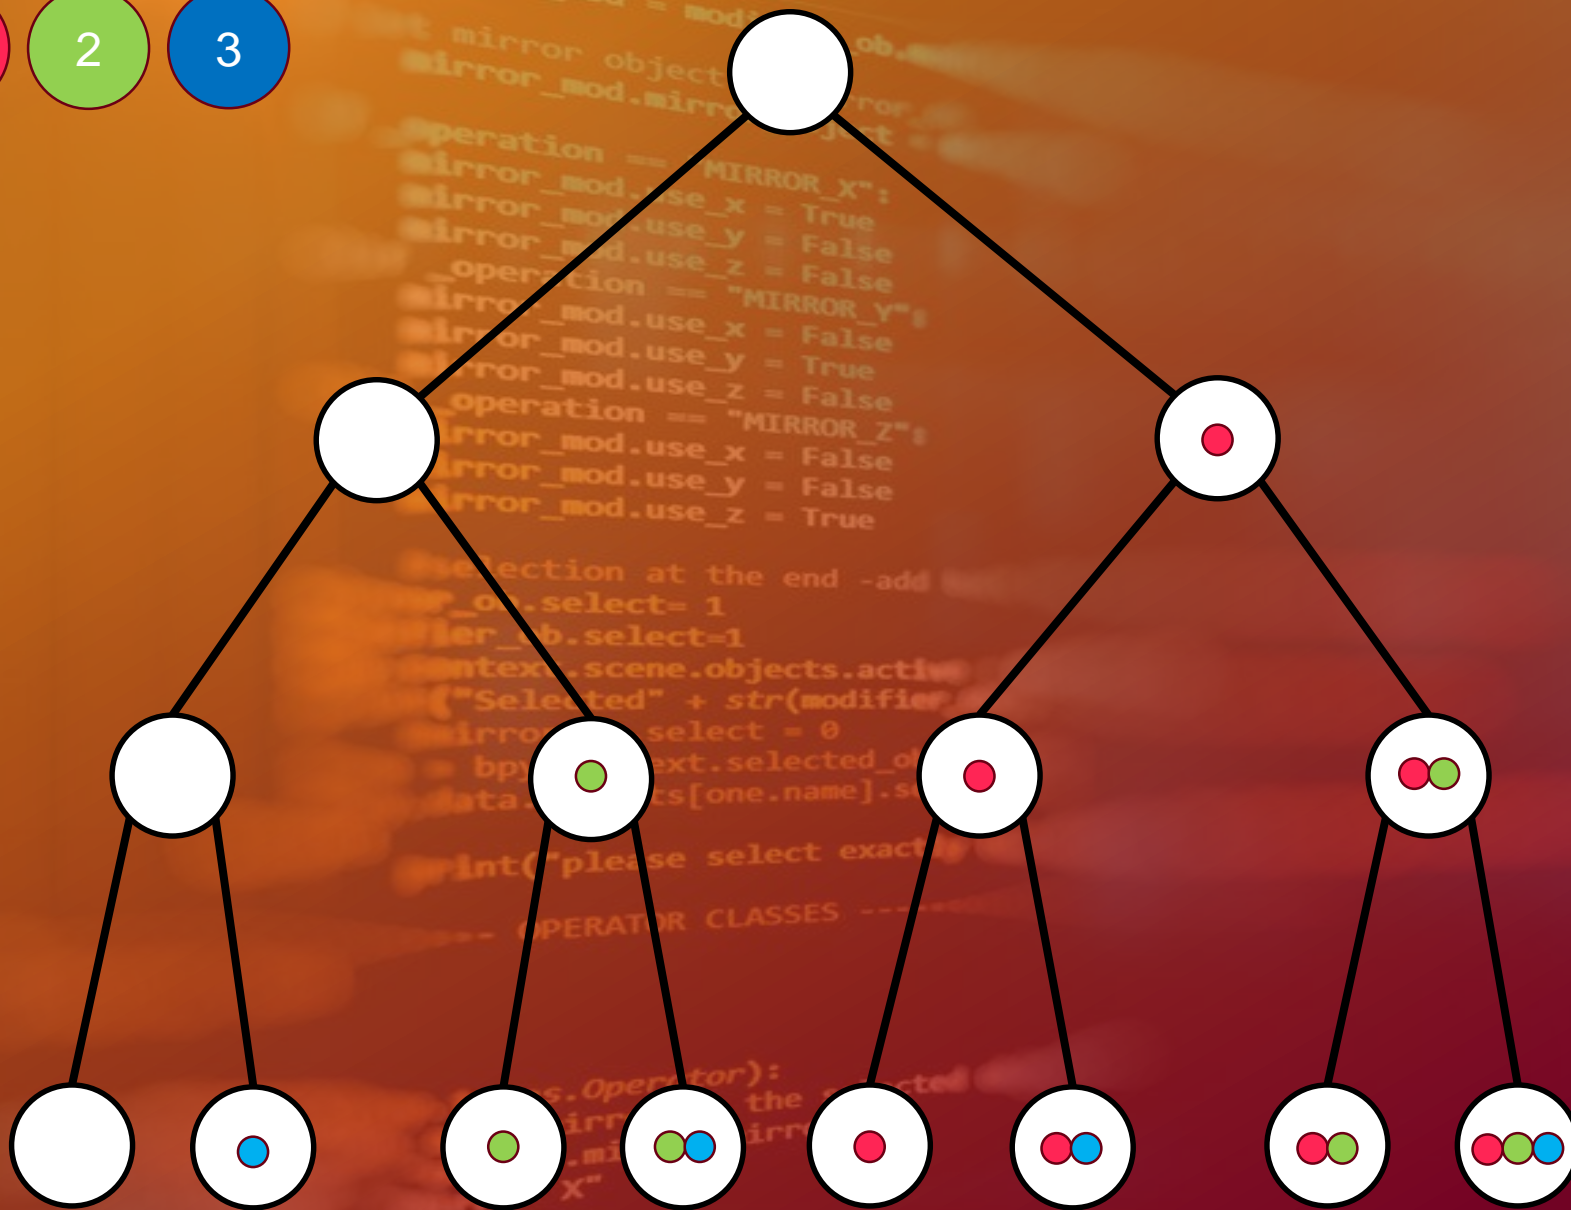


KNIGHTEC



Solutions = $n^m = 2^3 = 8$





- *Depth First Search*
- *Breadth First Search*
- *Morris in-order*
- ...

- *Depth First Search*
- *Breadth First Search*
- *Morris in-order*
- ...

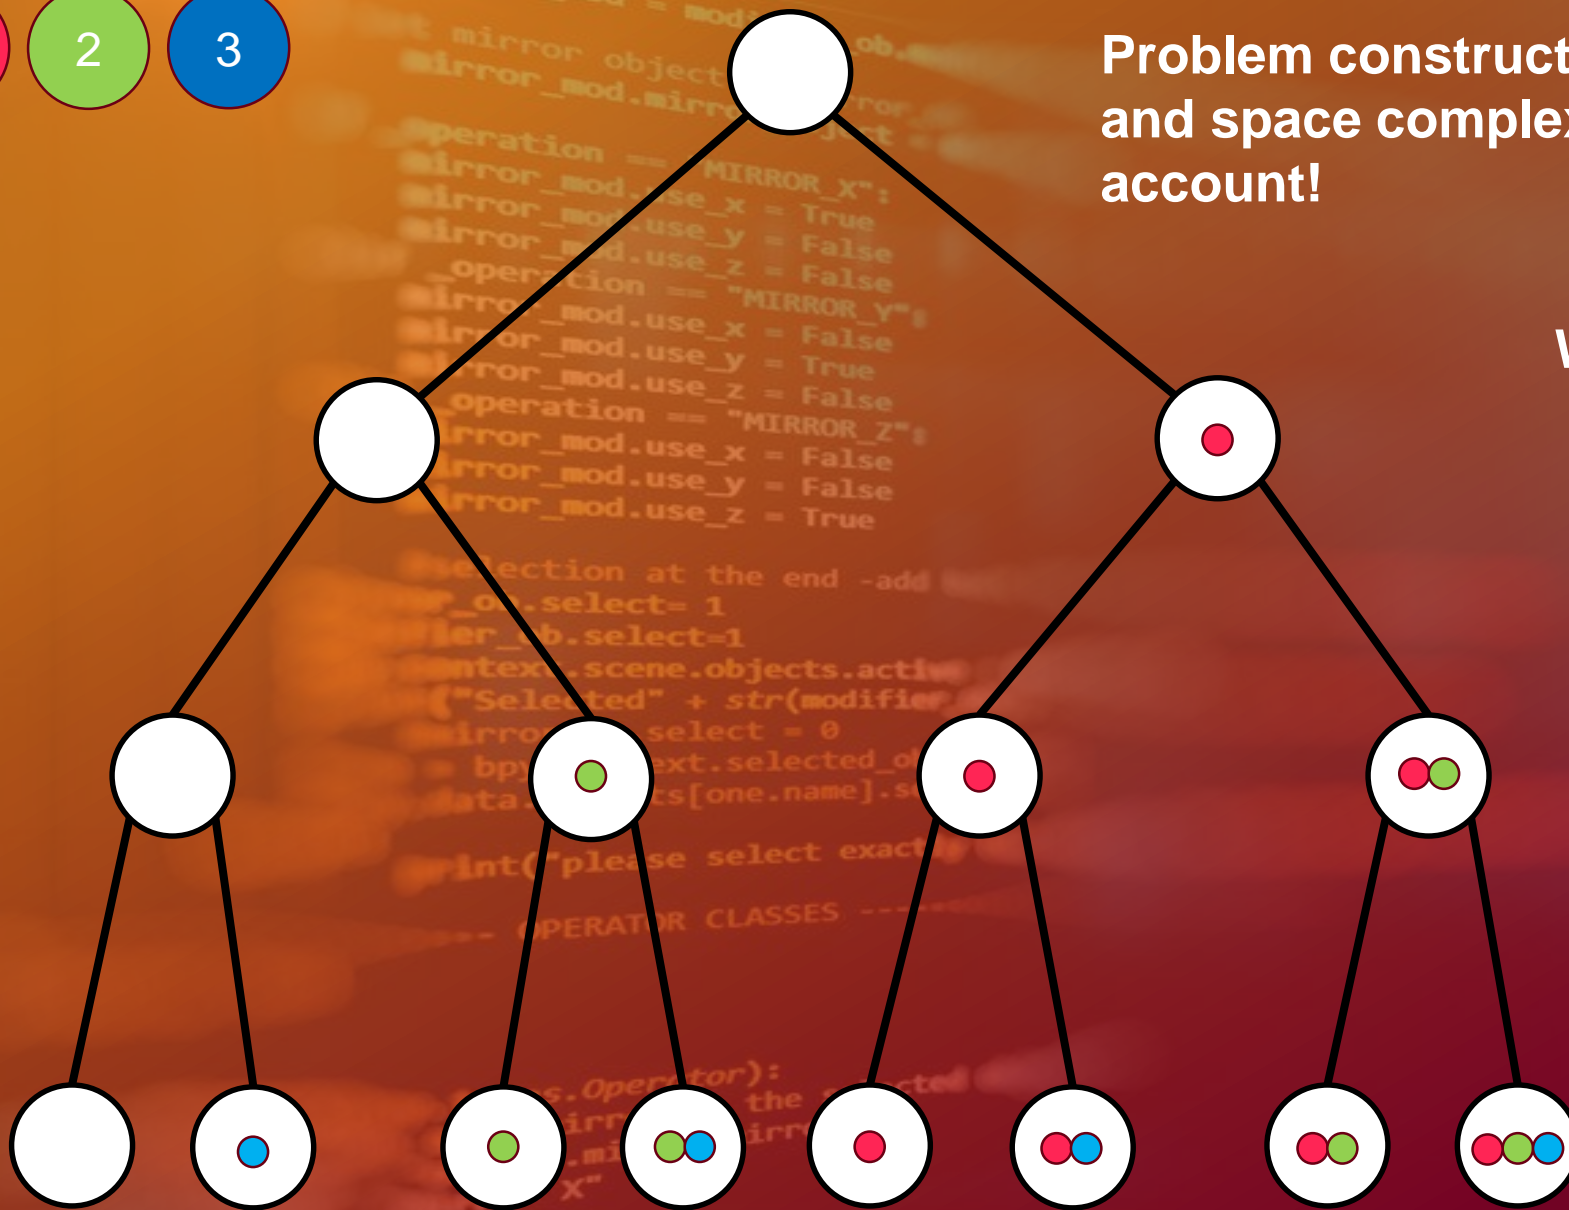
Demo

```
mirror_mod = modifier_ob.  
# Set mirror object to mirror  
mirror_mod.mirror_object =  
operation == "MIRROR_X":  
mirror_mod.use_x = True  
mirror_mod.use_y = False  
mirror_mod.use_z = False  
operation == "MIRROR_Y":  
mirror_mod.use_x = False  
mirror_mod.use_y = True  
mirror_mod.use_z = False  
operation == "MIRROR_Z":  
mirror_mod.use_x = False  
mirror_mod.use_y = False  
mirror_mod.use_z = True  
# selection at the end -add  
mirror_ob.select= 1  
modifier_ob.select=1  
context.scene.objects.active  
("Selected" + str(modifier_ob.name))  
mirror_ob.select = 0  
= bpy.context.selected_objects  
data.objects[one.name].select  
print("please select exactly one object")  
--- OPERATOR CLASSES ---  
bpy.types.Operator:  
X mirror to the selected  
object.mirror_mirror_x"  
mirror X"  
context):  
context.active_object is not None
```




Problem construction is important, and time- and space complexity must be taken into account!

Why is BFS worse than DFS?



A small, light-colored dog with dark spots and a red collar is looking towards a hand reaching out from the left. The background is a warm, orange-to-red gradient.

Cumulative rewarding and decision making

How can the ghosts know how to reach Pac-man?



Blinky the "chaser"



Pinky the "ambusher"



Inky the "random"




Clyde the "switcher"



Numerical matrix

1	2	2	2	2	2	2	2	2	3
4	0	0	0	0	0	0	0	0	8
4	0	X	0	0	X	X	0		8
4	0	X	0	0	X	0	0	0	8
4	0	X	0	0	0	0	X	0	8
4	0	X	0	0	X	X	X	0	8
4	0	0	0	0	0	0	0	0	8
5	6	6	6	6	6	6	6	6	7

Class-based representation




Reward-based construct,
based on current states

X	X	X	X	X	X	X	X	X	X
X	3		3	5	3	5	3	2	X
X	5	X	6	7	X	X	5	4	X
X	9		5	7	X	6	5	7	X
X	3	X	3	4	4	6	X	6	X
X	3	X	2	1	X	X	X	4	X
X	5	6	1	3	1		2	3	X
X	X	X	X	X	X	X	X	X	X

Ghost


- *Pac-man position*
- *Ghost position*




Square-class

- *Position [X, Y]*
- *Expanded/visited?*
- *Possible directions*
- *Actions performed from this square*

How many steps are required to reach Pac-man?





Birds-way-distance?








How do we solve this??? 0o

✖ = Maze wall

11	10	9	8	7	6	5	4	5	6
10	9		7	6	5	4	3	✖	5
9	8	7	6	5	4	3	✖	3	✖
8	7	6	5	4	3	✖	1	2	3
7	6	5	4	3	✖	1		✖	2
8	7	6	5	✖	3	2	1	✖	3
9	8	7	✖	✖	✖	✖	✖	✖	4
10	9	8	7	6	5	4	3	4	5

This is appears to be the best way... but we are stuck!

Best-first search, using A* and recursion

11	10	9	8	7	6	5	4	5	6
10	9		7	6	5	4	3		5
9	8	7	6	5	4	3		3	
8	7	6	5	4	3		1	2	3
7	6	5	4	3		1			2
8	7	6	5		3	2	1		3
9	8	7							4
10	9	8	7	6	5	4	3	4	5

$f(n) = g(n) + h(n)$

Break limit

Solution count limit

Number of steps limit

Random select

Demo

KNIGHTEC



Generation
1



Generation
6



Generation
17



Generation
921



Learning how to walk is never easy

Recommended reading:

“Q-Learning”

“Intelligent agents”



Q –learning A famous unsupervised learning algorithm, used among other things, for learning how to play games such as Super Mario, through “Trial and error”.

Did you notice the example running on the computer in this office?

Q-learning baby;)

GA - Evolution through "Survival of the fittest"

- Population of individuals – A breed or species to evolve
- Parents and offspring – The parents “compete” for the right to mate
- DNA/Chromosomes – Pass on the genetic inheritance to the offspring from both parents
- Mutation – Preserve genetic diversity (inbreeding control)
- Selection and Fitness function – Who will be allowed to breed?

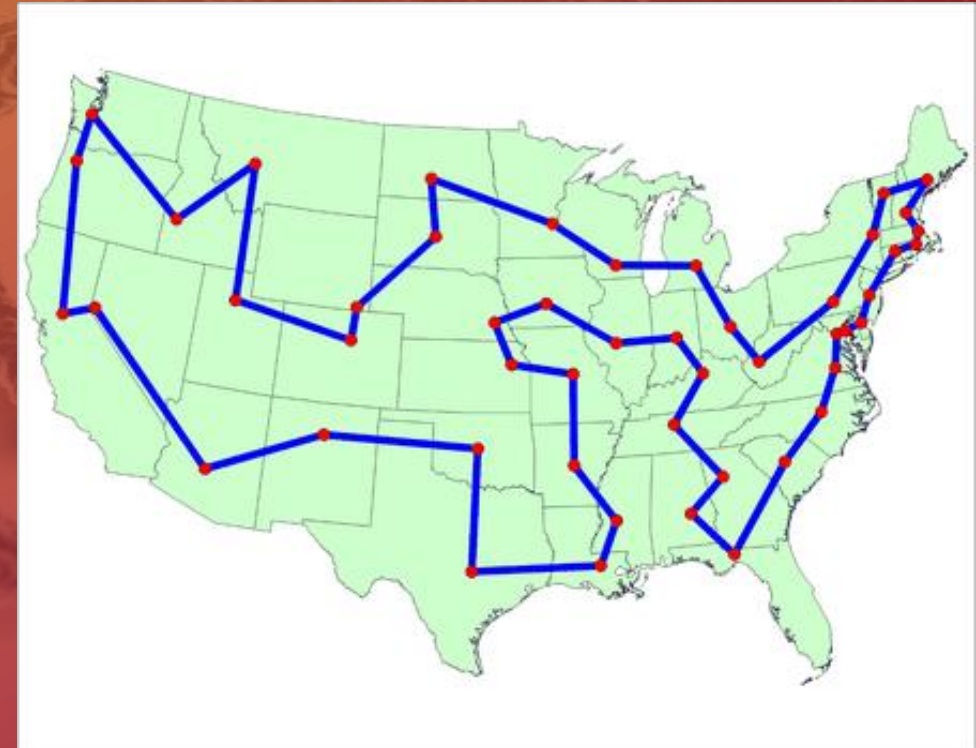
TSP – The Traveling salesman problem

Rules

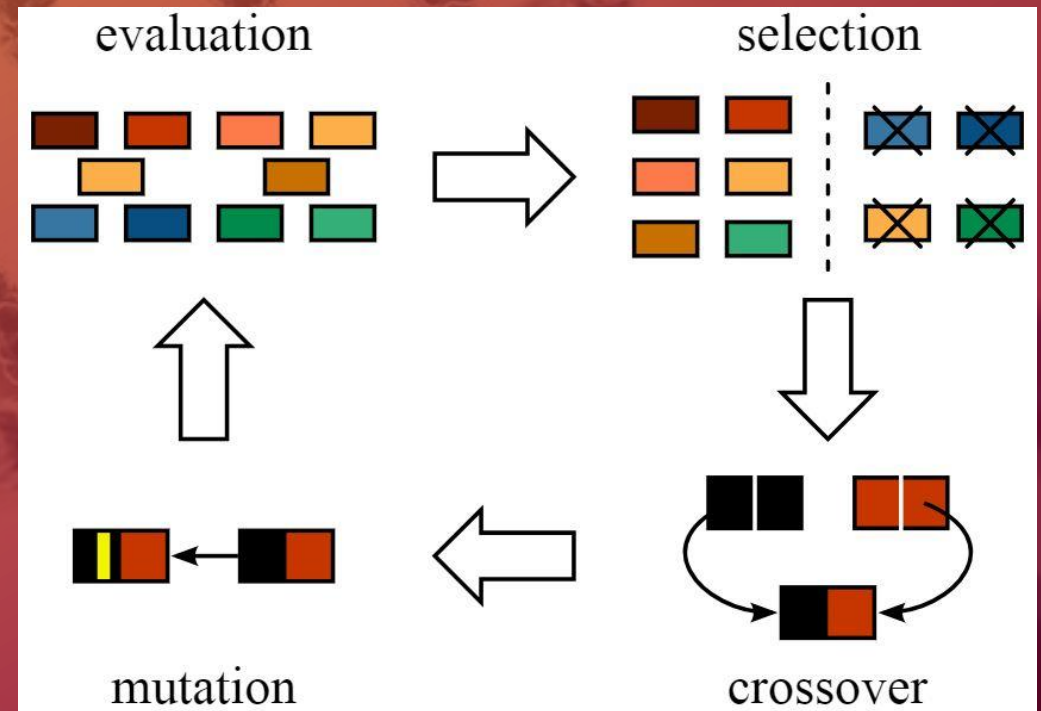
- Visit all cities once, and once only
- Minimize the total traveled distance
- Start and end the journey in the same city

Question – In terms of TSP, what corresponds to

1. the DNA/chromosomes?
2. a solution?
3. a population?
4. a parent?
5. an offspring?
6. the fitness?



1. Initiate the population randomly - Give each parent a random set of chromosomes (a list of cities/areas).
2. Select parents for breeding, based on the fitness function (best round-trip-distance)
3. Mate the set of selected parents together into couples
4. Pass DNA/genes to offspring, from both parents
5. Apply genetic mutations randomly on the new genetic composition
6. Repeat process on the new generation of individuals...







Demo

KNIGHTEC

Recommended further reading

- “Ant colony”
- “Flocking Birds”





Introduction to Neural Networks and classification

KNIGHTEC



Imagine we would build a machine that could identify fruit


What properties could be used by our machine, in order to distinguish fruits from each other?



Gather data

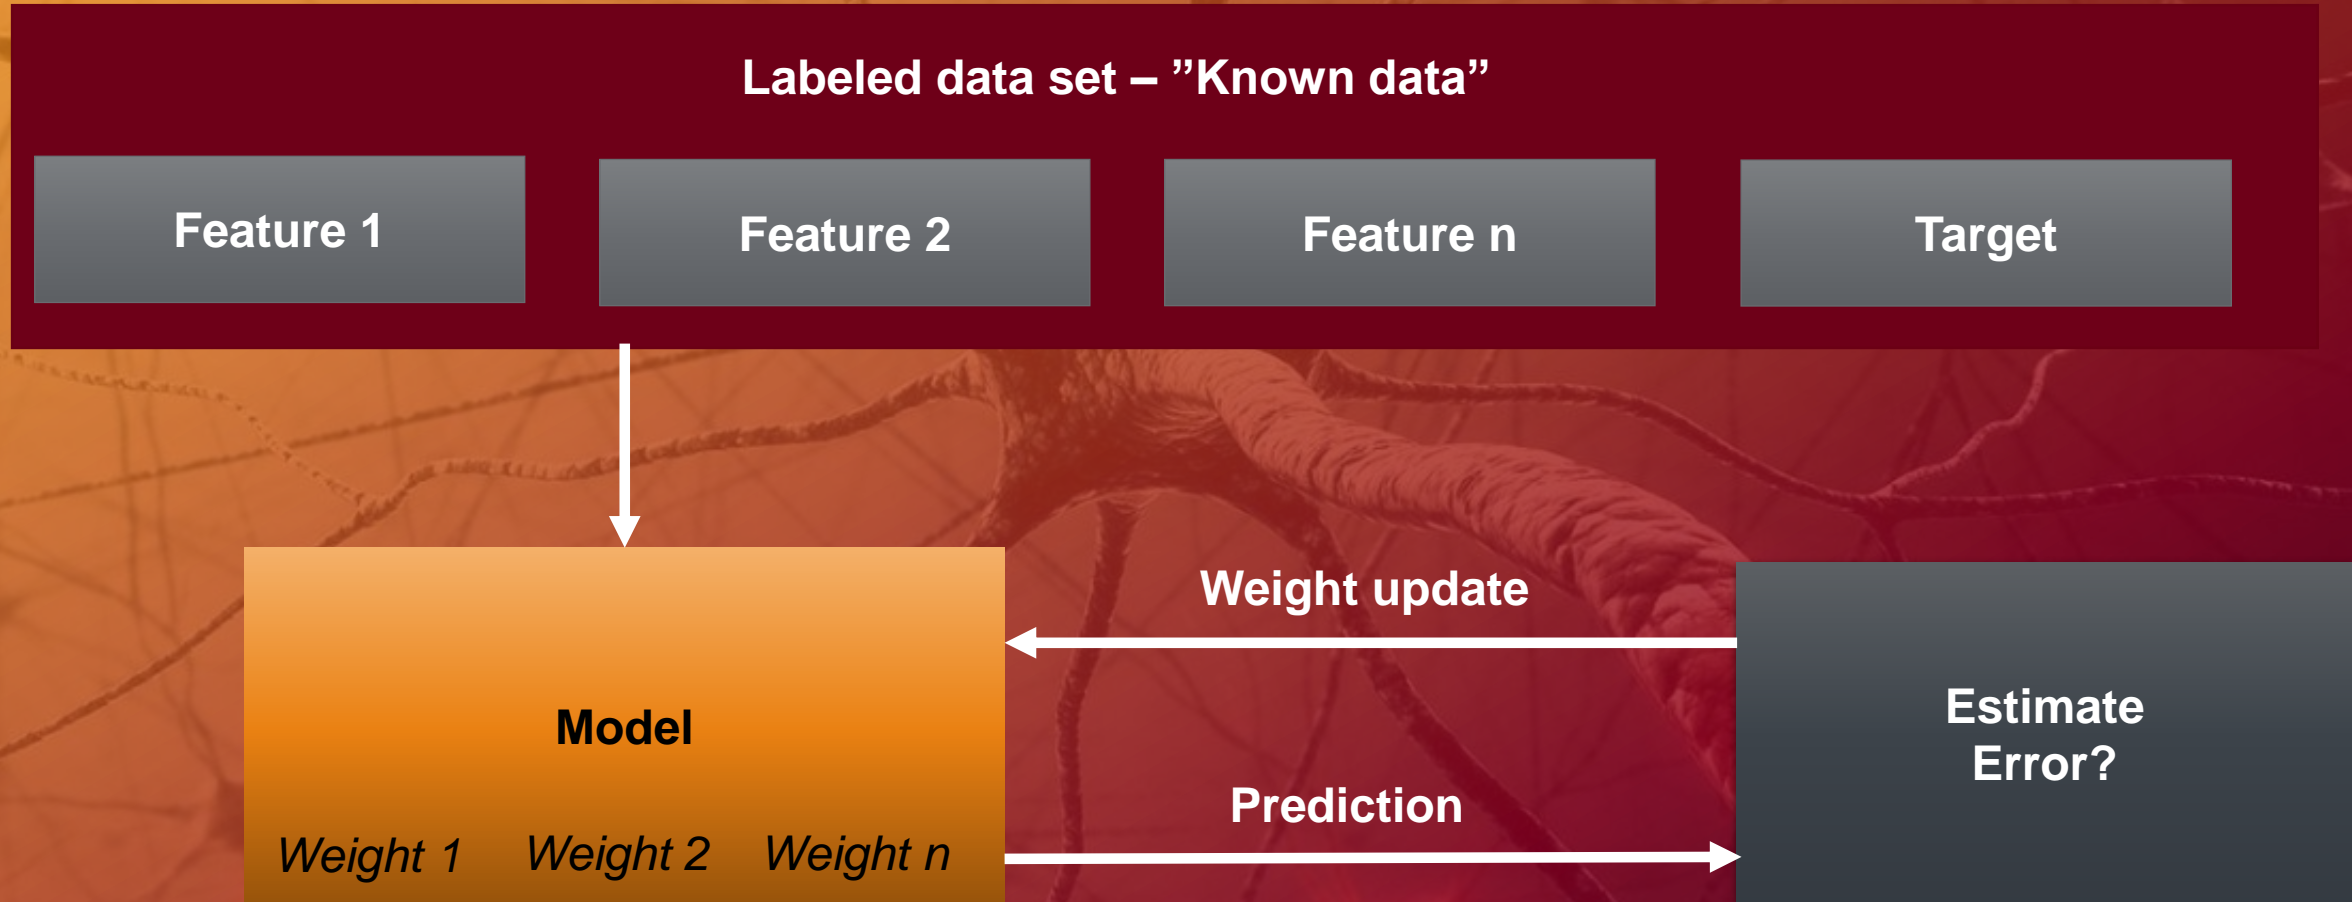


*"Known data set"/
"Labeled data"*

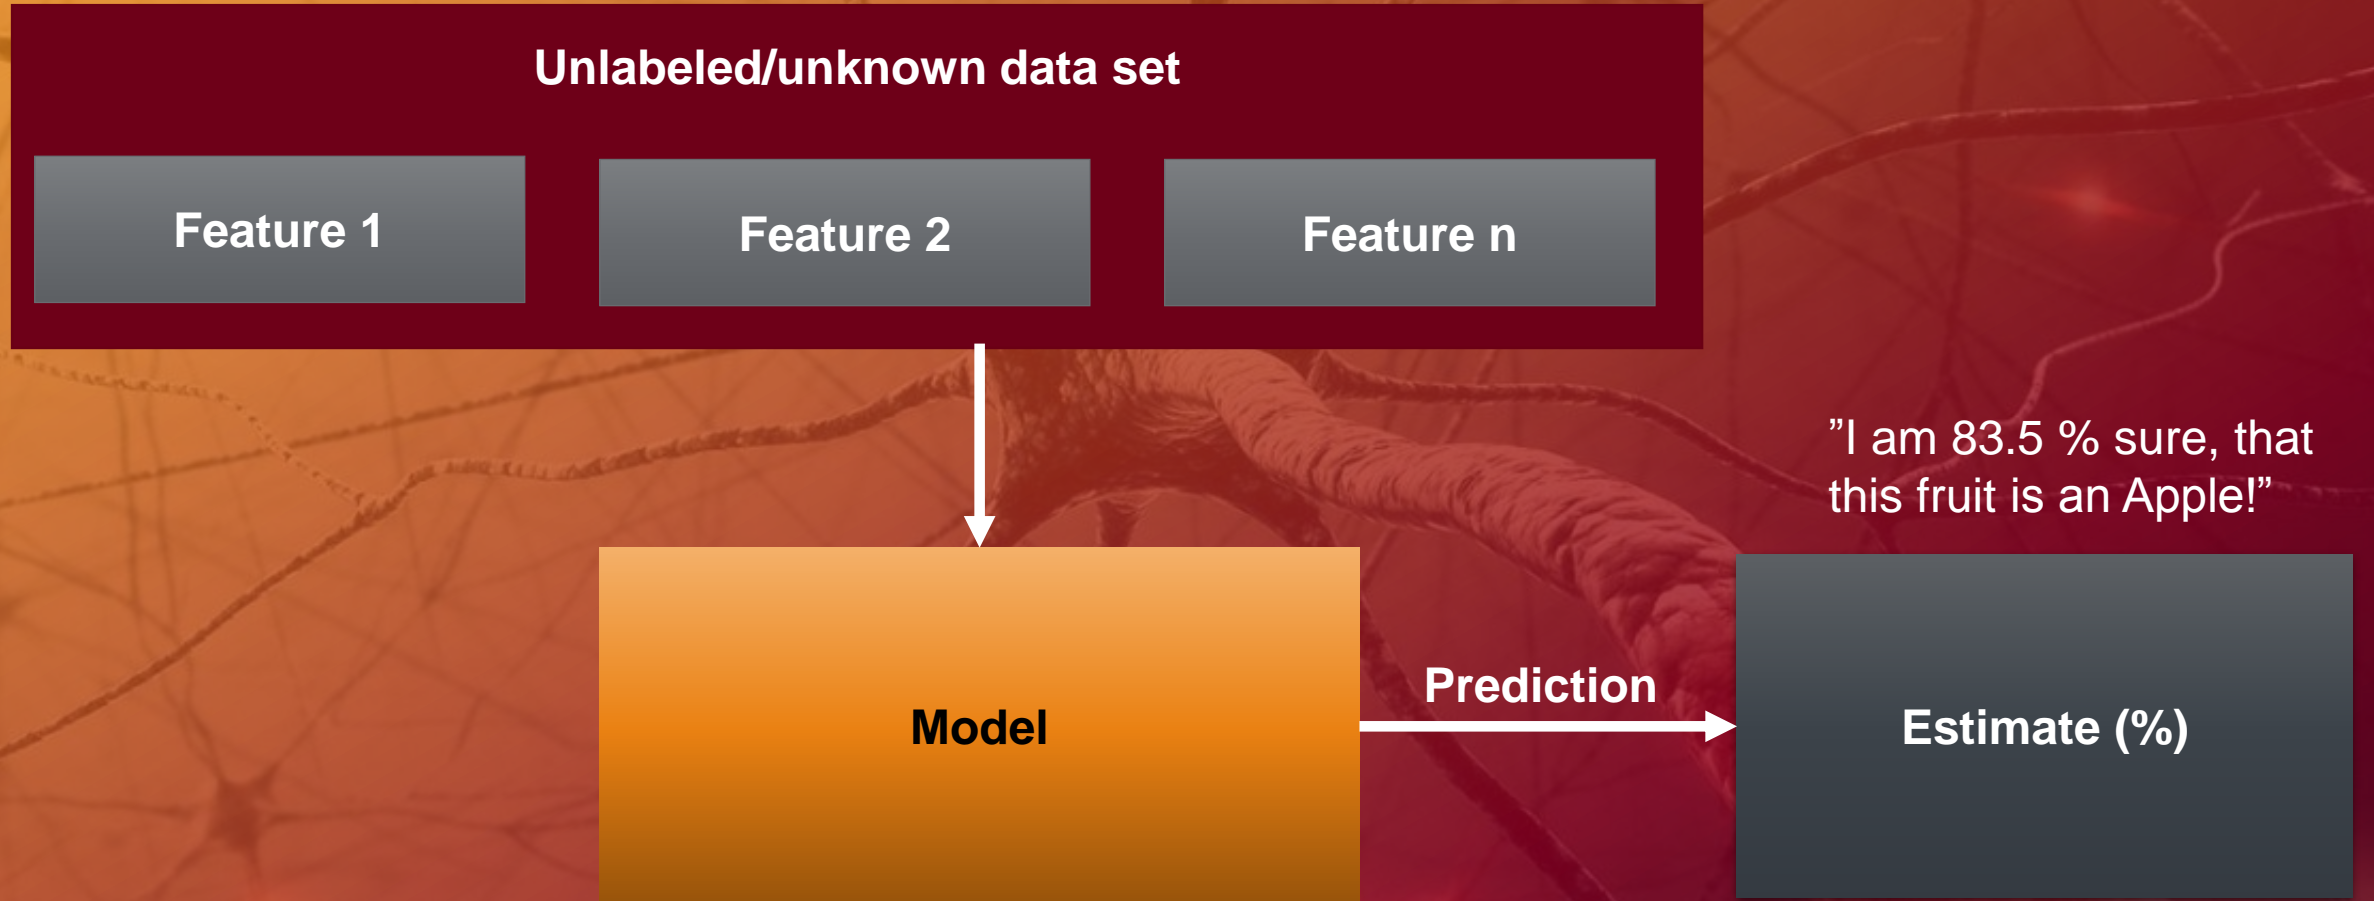


Width	Height	Weight	Color	Calories	Target
161.34	140.32	150.6	[249, 29, 42]	154	Apple
139.2	154.9	143.78	[80, 236, 51]	161	Pear
110-56	120.3	200.05	[228, 139, 63]	115	Orange
143.01	161.84	161.6	[235, 25, 62]	134	Apple
205.2	356.09	532.56	[190, 121, 42]	536	Pineapple
...
...
...
...
...

Training and model construction



Classification and prediction





Artificial Neuron

We don't have the required time to cover all the necessary math today, but I recommend you read about the different mathematical properties, and why they are used.

Don't panic now! Just try to accept that there is more to the mathematical construct of ANN, then the simplified explanations that follows.

Let's just try to get the basic idea =)



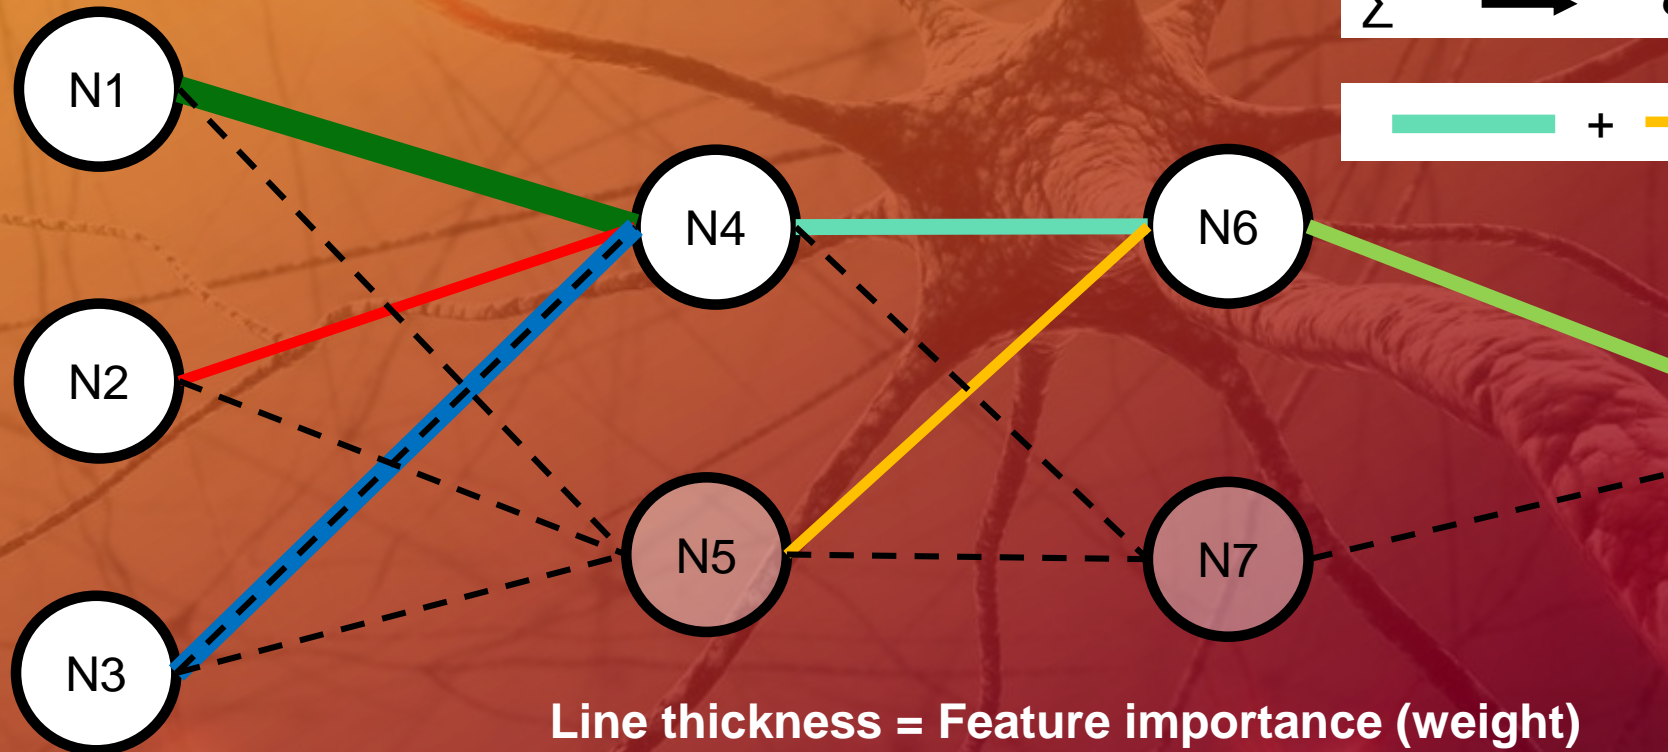
Neuron outputs in terms of RGB color

RGB[100, 220, 180] = 

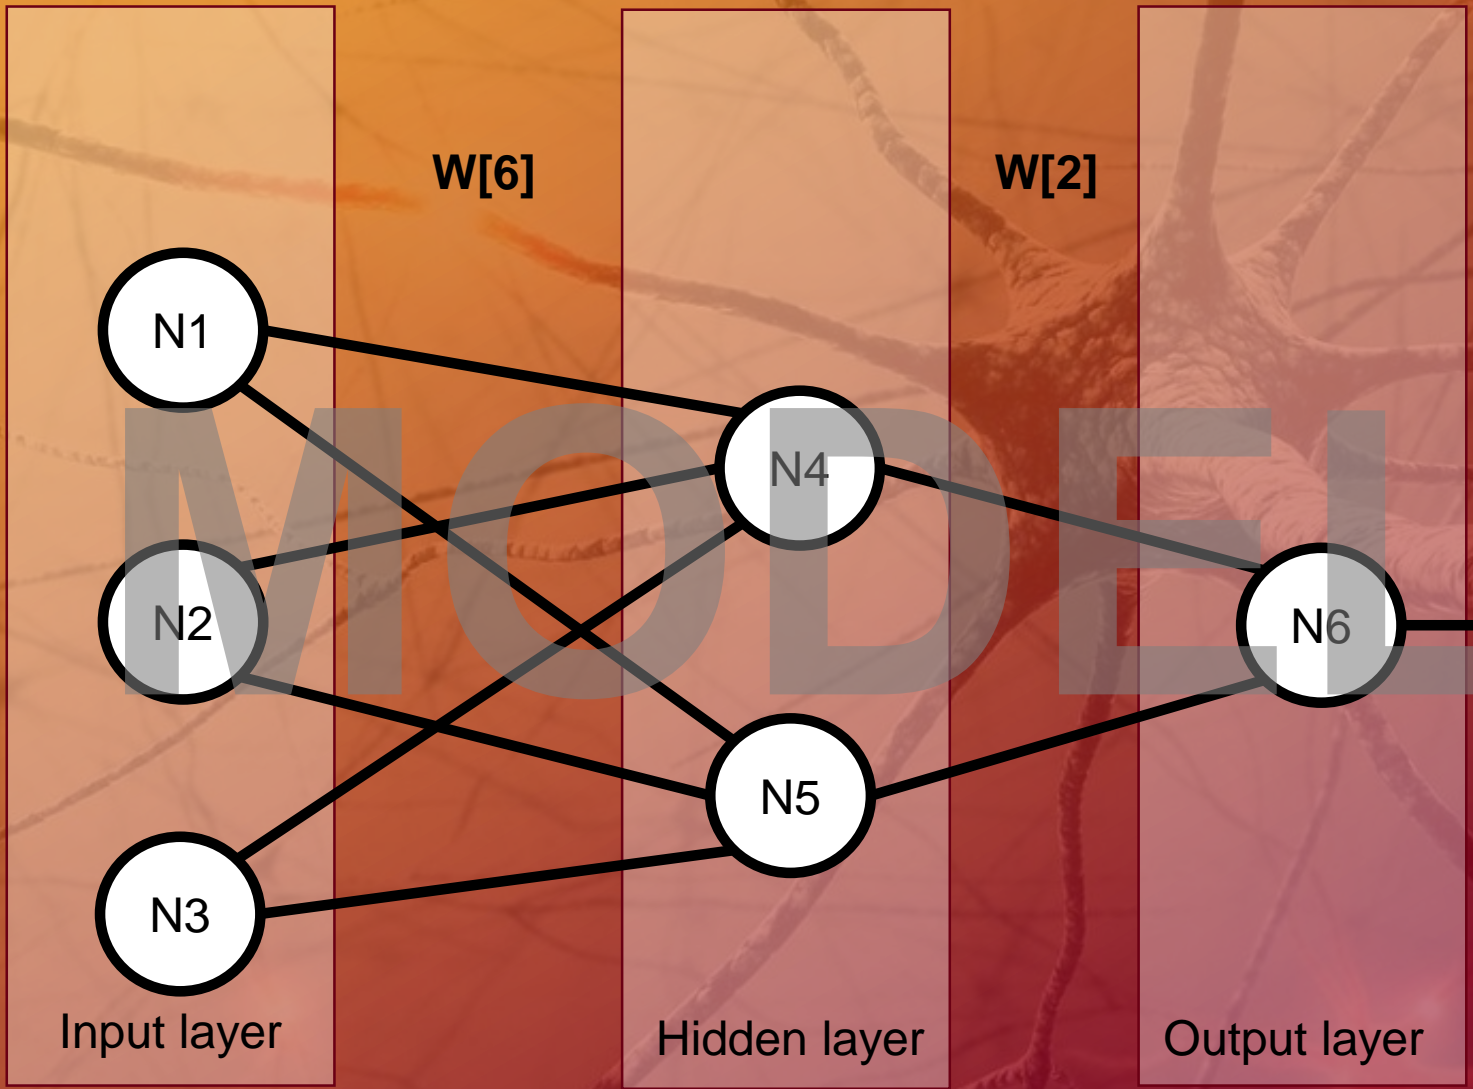
Feature	R	G	B
Weight (w)	100	220	180

$\Sigma \rightarrow \varsigma \rightarrow$ neuron output

 +  = 



Line thickness = Feature importance (weight)



- 1) Set initial weights
- 2) Calculate layers output, left -> right
- 3) Update layer weights, right->left, with respect to the estimation error.


Neuron output, and the “Sigmoid” activation function

“An acceptable mathematical representation of a biological neurons behavior.”

There is a lot more to it, but in short, we use this mathematical property to get an estimation in percentage (%), from the neuron output.

This could for example correspond to “how likely it is, that this particular fruit is an apple”.

Other activation functions exists, and depending on the desired neuron output characteristics, a different activation function could be used (step function, linear function, etc).



How do we change our weights?

Backpropagation and gradient decent

“Error function” - Difference between the estimate and the target value.

“Learning rate” – A numeric property on how “fast/steep” a correction will be (how aggressive the change of the updated weights will be).

The resulting ΔW is subject to both the estimation error, and the learning rate.

In the backpropagation, we change our weights with a calculated ΔW .

$$\text{New_W} = \text{prev_W} + \Delta W \text{ (simplified notation)}$$



Demo

KNIGHTEC

How could we use these tools on previous examples?

- The autonomous robots and the simulation environments?
- The Pac-man game?
- The TSP example?
- Titanic example?

Examples will be available in a GIT-repository for you to try out after the event. Experiment! =)

KNIGHTEC

THANK YOU FOR ATTENDING